

+



SAHYADRI
COLLEGE OF ENGINEERING & MANAGEMENT
An Autonomous Institution
MANGALURU

Department of Computer Science and Engineering
(Artificial Intelligence and Machine Learning)

DIGITAL IMAGE PROCESSING
(21AI71)


LAB RECORD

Submitted by,

Yagnik S Ram
4SF21CI057

Department of Computer Science and Engineering
(Artificial Intelligence and Machine Learning)

Marks



Head of The Department

Signature of the
Faculty-Incharge of the batch

INDEX – LABORATORY

Sl.no	Program	Date	Page No.
1	Splitting a digital image into four quadrants	20-09-2024	2
2	Applying edge detection and texture filters to extract low-level features.	27-09-2024	4
3	Enhancing and segmenting low contrast 2D images	04-10-2024	6
4	Image restoration method using spatial or frequency domain	18-10-2024	7
5	Compare edges by subtracting eroded and dilated images from the original.	12-11-2024	8
6	Image processing model using Computer Vision libraries	19-11-2024	10
7	Perform rotation, scaling, and translation on an image	26-11-2024	11

INDEX - ASSIGNMENT

Sl.no	Program	Page No.
1	Construct histogram for RGB image.	14
2	Use linear, gamma and logarithmic transformations to an image.	16
3	Construct the histogram for a grayscale image.	17
4	Convert a given image to a negative image.	19
5	Convert RGB image to grayscale and display R, G, B channels.	20
6	Pixelate, blur, and add noise to an image.	22
7	Enhance image contrast using histogram equalization.	24

LAB PROGRAMS

1. Write a Program to read a digital image. Split and display image into 4 quadrants, up, down, right and left

```
# pip install opencv-python
import cv2

Image_path = 'LAB_1/Miguel.jpeg'
image = cv2.imread(image_path)
if image is None:
    print("Error: Unable to load image.")
    exit()

print(image.shape)
height, width, _ = image.shape

center_y, center_x = height // 2, width // 2


top_left = image[0:center_y, 0:center_x]
top_right = image[0:center_y, center_x:]
bottom_left = image[center_y:, 0:center_x]
bottom_right = image[center_y:, center_x:]

cv2.imshow('Top Left', top_left)
cv2.imshow('Top Right', top_right)
cv2.imshow('Bottom Left', bottom_left)
cv2.imshow('Bottom Right', bottom_right)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

INPUT IMAGE



OUTPUT



2. Read an image and extract and display low-level features such as edges, textures using filtering techniques

```
import cv2

# Function to apply Gaussian Filter
def apply_gaussian_filter(image, kernel_size=5, sigma=1.0):
    # Apply Gaussian Blur
    gaussian_filtered = cv2.GaussianBlur(image, (kernel_size, kernel_size), sigma)
    return gaussian_filtered

# Function to apply Median Filter
def apply_median_filter(image, kernel_size=13):
    # Apply Median Filter
    median_filtered = cv2.medianBlur(image, kernel_size)
    return median_filtered

# Main function
def main():
    # Hard-coded image path
    image_path = 'LAB_2/flower.jpg' # Change this to your image path

    # Read the image
    image = cv2.imread(image_path)

    if image is None:
        print("ERROR! Could not read the image.")
    return

    # Resize the image to a smaller size (e.g., 40% of the original)
    new_width = int(image.shape[1] * 0.4)
    new_height = int(image.shape[0] * 0.4)
    resized_image = cv2.resize(image, (new_width, new_height))

    # Apply Gaussian Filter
    gaussian_result = apply_gaussian_filter(resized_image, kernel_size=5, sigma=1.0)

    # Apply Median Filter
    median_result = apply_median_filter(resized_image, kernel_size=5)

    # Display the results
    cv2.imshow('Original Image (Resized)', resized_image)
    cv2.imshow('Gaussian Filtered Image', gaussian_result)
    cv2.imshow('Median Filtered Image', median_result)
```



SAHYADRI
COLLEGE OF ENGINEERING & MANAGEMENT
An Autonomous Institution
MANGALURU

Wait until a key is pressed, then close all windows

```
cv2.waitKey(0)
```

```
cv2.destroyAllWindows()
```

```
if __name__ == "__main__":  
    main()
```

OUTPUT



Original Image



Gaussian Filtered



Median Filtered

3. Demonstrate enhancing and segmenting low contrast 2D images

```
import cv2
import numpy as np

# Load the image
image = cv2.imread('LAB_3\low_contrast_image.png', cv2.IMREAD_GRAYSCALE)

# Enhance the image using CLAHE
clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8, 8))
enhanced_image = clahe.apply(image)

# Segment the image using Otsu's thresholding
_, segmented_image = cv2.threshold(enhanced_image, 0, 255, cv2.THRESH_BINARY +
cv2.THRESH_OTSU)

# Display the results
cv2.imshow('Original Image', image)
cv2.imshow('Enhanced Image', enhanced_image)
cv2.imshow('Segmented Image', segmented_image)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

OUTPUT



Original Image



Enhanced Image



Segmented Image



4. Demonstrate image restoration using spatial or frequency domain

```
import cv2
import numpy as np
import matplotlib.pyplot as plt

# Load the image
image = cv2.imread('LAB_4/noisy_dog.jpg', 0)

# Spatial Domain Example: Gaussian Filter
spatial_filtered = cv2.GaussianBlur(image, (5, 5), 0)

# Frequency Domain Example
# Step 1: Fourier Transform
f = np.fft.fft2(image)
fshift = np.fft.fftshift(f)

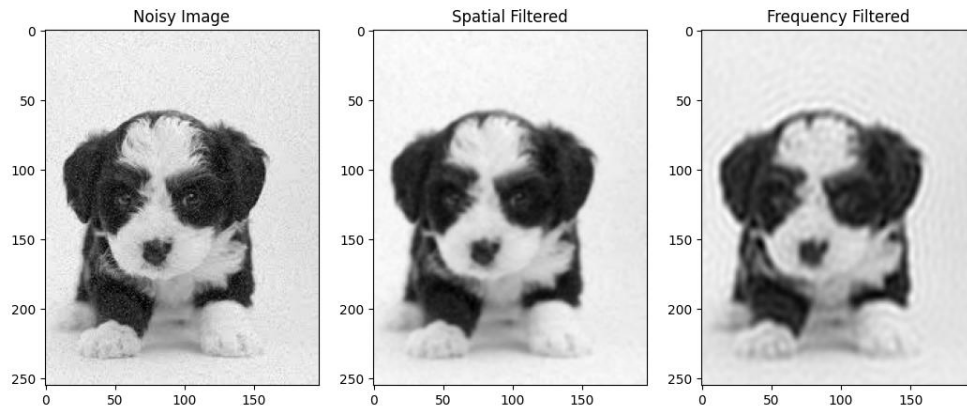
# Step 2: Create a mask (low-pass filter)
rows, cols = image.shape
crow, ccol = rows // 2, cols // 2
mask = np.zeros((rows, cols), np.uint8)
r = 30 # Radius of the mask
cv2.circle(mask, (ccol, crow), r, 1, thickness=-1)

# Step 3: Apply the mask
fshift_filtered = fshift * mask

# Step 4: Inverse Fourier Transform
f_ishift = np.fft.ifftshift(fshift_filtered)
image_restored = np.fft.ifft2(f_ishift)
image_restored = np.abs(image_restored)

# Display the results
plt.figure(figsize=(12, 6))
plt.subplot(1, 3, 1), plt.imshow(image, cmap='gray'), plt.title('Noisy Image')
plt.subplot(1, 3, 2), plt.imshow(spatial_filtered, cmap='gray'), plt.title('Spatial Filtered')
plt.subplot(1, 3, 3), plt.imshow(image_restored, cmap='gray'), plt.title('Frequency Filtered')
plt.show()
```


OUTPUT



5. Read an image, first apply erosion to the image and then subtract the result from the original. Demonstrate the difference in the edge image if you use dilation instead of erosion.

```
import cv2
import numpy as np
import matplotlib.pyplot as plt

# Read the image
image = cv2.imread('LAB_5/anime.jpg', cv2.IMREAD_GRAYSCALE)

# Define the kernel for erosion and dilation
kernel = np.ones((5, 5), np.uint8)

# Apply erosion
eroded_image = cv2.erode(image, kernel, iterations=1)

# Subtract the eroded image from the original
subtracted_image = cv2.subtract(image, eroded_image)

# Apply dilation
dilated_image = cv2.dilate(image, kernel, iterations=1)

# Subtract the dilated image from the original

# Display the results
```

```
plt.figure(figsize=(10, 8))

plt.subplot(2, 2, 1)
plt.title('Original Image')
plt.imshow(image, cmap='gray')


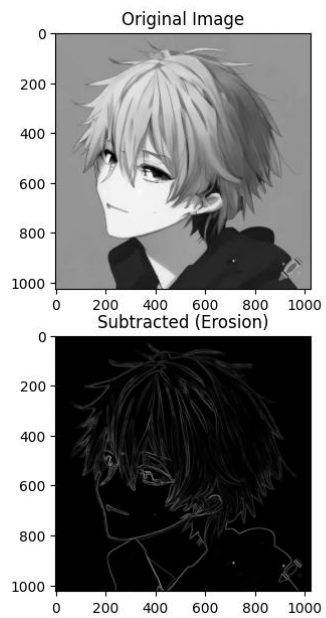
plt.subplot(2, 2, 2)
plt.title('Eroded Image')
plt.imshow(eroded_image, cmap='gray')

plt.subplot(2, 2, 3)
plt.title('Subtracted (Erosion)')
plt.imshow(subtracted_image, cmap='gray')

plt.subplot(2, 2, 4)
plt.title('Subtracted (Dilation)')
plt.imshow(dilated_image, cmap='gray')

plt.tight_layout()
```

OUTPUT



6. Implement image processing model using Computer Vision libraries

(TensorFlow, Keras)

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras import layers, models from tensorflow.keras.datasets import mnist
```

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train.reshape((x_train.shape[0], 28, 28, 1)).astype('float32') / 255
x_test = x_test.reshape((x_test.shape[0], 28, 28, 1)).astype('float32') / 255 y_train =
tf.keras.utils.to_categorical(y_train, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)
```

```
model = models.Sequential([
layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation='relu'), layers.Flatten(),
layers.Dense(64, activation='relu'), layers.Dense(10, activation='softmax')
])
```


```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, epochs=10, batch_size=128, validation_split=0.2)
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_acc:.4f}')
```

```
num_results = 5
predictions = model.predict(x_test)
```

```
plt.figure(figsize=(12, 6)) for i in range(num_results):
plt.subplot(1, num_results, i + 1) plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
plt.title(f'Pred: {np.argmax(predictions[i])}, True: {np.argmax(y_test[i])}') plt.axis('off')
plt.tight_layout() plt.show()
```

OUTPUT

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 — 14s 1us/step
Epoch 1/10
375/375 — 12s 25ms/step - accuracy: 0.7982 - loss: 0.6498 - val_accuracy: 0.9718 - val_loss: 0.0937
Epoch 2/10
375/375 — 9s 24ms/step - accuracy: 0.9748 - loss: 0.0808 - val_accuracy: 0.9818 - val_loss: 0.0588
Epoch 3/10
375/375 — 9s 25ms/step - accuracy: 0.9858 - loss: 0.0485 - val_accuracy: 0.9869 - val_loss: 0.0458
Epoch 4/10
375/375 — 9s 24ms/step - accuracy: 0.9899 - loss: 0.0327 - val_accuracy: 0.9857 - val_loss: 0.0508
Epoch 5/10
375/375 — 9s 24ms/step - accuracy: 0.9921 - loss: 0.0269 - val_accuracy: 0.9875 - val_loss: 0.0430
Epoch 6/10
375/375 — 9s 24ms/step - accuracy: 0.9932 - loss: 0.0202 - val_accuracy: 0.9883 - val_loss: 0.0397
Epoch 7/10
375/375 — 10s 25ms/step - accuracy: 0.9938 - loss: 0.0189 - val_accuracy: 0.9887 - val_loss: 0.0398
Epoch 8/10
375/375 — 10s 26ms/step - accuracy: 0.9954 - loss: 0.0140 - val_accuracy: 0.9878 - val_loss: 0.0418
Epoch 9/10
375/375 — 9s 25ms/step - accuracy: 0.9958 - loss: 0.0140 - val_accuracy: 0.9902 - val_loss: 0.0392
Epoch 10/10
375/375 — 9s 25ms/step - accuracy: 0.9965 - loss: 0.0115 - val_accuracy: 0.9899 - val_loss: 0.0406
313/313 — 2s 5ms/step - accuracy: 0.9868 - loss: 0.0390
Test accuracy: 0.9909
313/313 — 2s 5ms/step
```



7. Write a program to show rotation, scaling, and translation of an image

```
import cv2
import numpy as np

# Load the image
image = cv2.imread('image.jpg')

# Get the dimensions of the image
(h, w) = image.shape[:2]

# Define the center of the image
center = (w // 2, h // 2)

# Rotation
angle = 45 # Rotate by 45 degrees
scale = 1.0 # No scaling during rotation
rotation_matrix = cv2.getRotationMatrix2D(center, angle, scale)
rotated_image = cv2.warpAffine(image, rotation_matrix, (w, h))

# Scaling
scale_x = 1.5 # Scale by 1.5 times along the x-axis
```



```
scale_y = 1.5 # Scale by 1.5 times along the y-axis
scaled_image = cv2.resize(image, None, fx=scale_x, fy=scale_y,
interpolation=cv2.INTER_LINEAR)

# Translation
tx = 100 # Translate by 100 pixels along the x-axis
ty = 50 # Translate by 50 pixels along the y-axis
translation_matrix = np.float32([[1, 0, tx], [0, 1, ty]])
translated_image = cv2.warpAffine(image, translation_matrix, (w, h))

# Display the results
cv2.imshow('Original Image', image)
cv2.imshow('Rotated Image', rotated_image)
cv2.imshow('Scaled Image', scaled_image)
cv2.imshow('Translated Image', translated_image)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

OUTPUT



Original Image



Rotated Image



Scaled Image



Translated Image



ASSIGNMENT PROBLEMS

1. Write a program to construct the histogram of an RGB image.

```
import cv2

import numpy as np

import matplotlib.pyplot as plt

# Load an HD image with RGB layers (replace 'image.jpg' with
your image path)

image = cv2.imread('CLASS_PROBLEMS/Images/Flower.jpeg')

image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) #
Convert to RGB

# Apply 2x2 average filter

kernel_2x2 = np.ones((2, 2), np.float32) / 4

filtered_2x2 = cv2.filter2D(image_rgb, -1, kernel_2x2)

# Apply 3x3 average filter

kernel_3x3 = np.ones((3, 3), np.float32) / 9

filtered_3x3 = cv2.filter2D(image_rgb, -1, kernel_3x3)

# Plot the original and filtered images

plt.figure(figsize=(15, 10))

plt.subplot(1, 3, 1)
```

```
plt.imshow(image_rgb)
```

```
plt.title("Original Image")
```

```
plt.axis("off")
```

```
plt.subplot(1, 3, 2)
```

```
plt.imshow(filtered_2x2)
```

```
plt.title("2x2 Average Filter")
```

```
plt.axis("off")
```

```
plt.subplot(1, 3, 3)
```


```
plt.imshow(filtered_3x3)
```

```
plt.title("3x3 Average Filter")
```

```
plt.axis("off")
```

```
plt.show()
```

OUTPUT





2. Write a code to apply linear, gamma, and logarithmic transformations to an image.

```
import cv2
import numpy as np

# Load the image
image = cv2.imread('image.jpg', cv2.IMREAD_GRAYSCALE)

# Linear Transformation
alpha = 1.5 # Simple contrast control
beta = 50   # Simple brightness control
linear_transformed = cv2.convertScaleAbs(image, alpha=alpha, beta=beta)

# Gamma Transformation
gamma = 2.0
gamma_corrected = np.array(255 * (image / 255) ** gamma, dtype='uint8')

# Logarithmic Transformation
c = 255 / np.log(1 + np.max(image))
log_transformed = c * (np.log(image + 1))
log_transformed = np.array(log_transformed, dtype='uint8')

# Display the results
cv2.imshow('Original Image', image)
cv2.imshow('Linear Transformed Image', linear_transformed)
cv2.imshow('Gamma Corrected Image', gamma_corrected)
cv2.imshow('Logarithmic Transformed Image', log_transformed)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

OUTPUT



Original Image



Linear Transformed Image



Gamma Corrected image



Logarithmic Transformed Image

3. Write a program to construct the histogram for a grayscale image.

```
import cv2
import numpy as np
import matplotlib.pyplot as plt

# Load the image using OpenCV
image =
cv2.imread('CLASS_PROBLEMS/Images/Flower.jpeg',
```



```
cv2.IMREAD_COLOR)
```

```
# Convert the image to RGB (OpenCV  
loads images in BGR format by  
default)
```

```
image_rgb = cv2.cvtColor(image,  
cv2.COLOR_BGR2RGB)
```

```
# Create a figure for the bar plot  
plt.figure(figsize=(10, 6))
```

```
# Initialize the number of bins (256 for  
pixel values 0-255)  
bins = 256
```

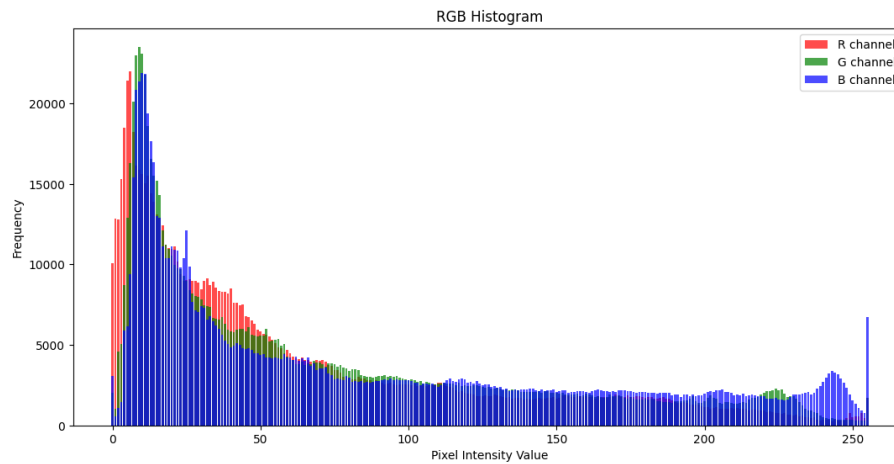
```
# Create a histogram for each color  
channel and plot bars  
colors = ('r', 'g', 'b') # Red, Green,  
Blue channels  
for i, color in enumerate(colors):  
    hist = cv2.calcHist([image_rgb], [i],  
None, [bins], [0, 256])  
    hist = hist.flatten() # Flatten the  
histogram array
```

```
# Create bar plot for each channel  
plt.bar(np.arange(bins), hist,  
color=color, alpha=0.7,  
label=f'{color.upper()} channel')
```

```
# Add labels and title  
plt.title('RGB Histogram')  
plt.xlabel('Pixel Intensity Value')  
plt.ylabel('Frequency')  
plt.legend()
```

```
# Display the histogram  
plt.show()
```

OUTPUT



4. Write Python code to convert a given image to a negative image.

```
import cv2
```

```
import numpy as np
```

```
# Load the image
```

```
image = cv2.imread('image.jpg')
```

```
# Convert the image to a negative image
```

```
negative_image = cv2.bitwise_not(image)
```

```
# Display the results
```

```
cv2.imshow('Original Image', image)
```

```
cv2.imshow('Negative Image', negative_image)
```

```
cv2.waitKey(0)
```

```
cv2.destroyAllWindows()
```

OUTPUT



Original Image



Negative Image

5. Write a Python code to convert an RGB image to grayscale and display the R, G, and B components of the image.

```
import cv2
import matplotlib.pyplot as plt
import numpy as np

def display_images(images, titles):
    plt.figure(figsize=(15, 5))
    for i, (image, title) in enumerate(zip(images, titles)):
        plt.subplot(1, len(images), i + 1)
        if len(image.shape) == 2: # Grayscale image
            plt.imshow(image, cmap='gray')
        else:
            plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
        plt.title(title)
        plt.axis('off')

plt.show() image_path = "1.jpg"
image = cv2.imread(image_path)

if image is None:
    print("Error: Could not read the image. Please check the file path.")
else:
    B, G, R = cv2.split(image)


    R_colored = cv2.merge([np.zeros_like(B), np.zeros_like(G), R]) # Red channel
    G_colored = cv2.merge([np.zeros_like(B), G, np.zeros_like(R)]) # Green channel
```

```
B_colored = cv2.merge([B, np.zeros_like(G), np.zeros_like(R)]) # Blue channel

grayscale_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

display_images(
    [image, grayscale_image, R_colored, G_colored, B_colored],
    ["Original Image", "Grayscale Image", "Red Component", "Green Component", "Blue Component"]
)
```

OUTPUT



Original Image



Grayscale Image



Red Component



Green Component



Blue Component

6. Write a Python code to pixelate, blur, and add noise to an image.

```
import cv2

import numpy as np

import matplotlib.pyplot as plt

# Load an HD image

image = cv2.imread('CLASS_PROBLEMS/Images/Sky.jpg')

image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) #

Convert to RGB format

image_rgb = cv2.resize(image_rgb, (720, 1080))

# Function to add Gaussian noise

def add_gaussian_noise(image, mean=0, stddev=25):

    # Generate Gaussian noise

    noise = np.random.normal(mean, stddev, image.shape)
```

```

noise = noise.reshape(image.shape)

# Add noise to the image and clip values to maintain valid pixel
range

noisy_image = image + noise

noisy_image = np.clip(noisy_image, 0, 255) # Keep pixel values
within [0, 255]

return noisy_image.astype(np.uint8)

# Spoil the image by adding Gaussian noise

spoiled_image = add_gaussian_noise(image_rgb, mean=0,
stddev=30)

# Display the original and spoiled images

plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)

plt.imshow(image_rgb)

plt.title("Original Image")

plt.axis("off")

plt.subplot(1, 2, 2)

plt.imshow(spoiled_image)

plt.title("Spoiled Image with Gaussian Noise")

plt.axis("off")

plt.show()

```


OUTPUT

Original Image



Spoiled Image with Gaussian Noise



7. Write a program to enhance contrast of an image using histogram equalisation. Display the results with and without equalisation for comparison.

```
import cv2
import numpy as np

# Load the image
image = cv2.imread('image.jpg', cv2.IMREAD_GRAYSCALE)

# Apply histogram equalization
equalized_image = cv2.equalizeHist(image)

# Display the results
cv2.imshow('Original Image', image)
cv2.imshow('Equalized Image', equalized_image)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

OUTPUT



Original Image



Equalized image